

Rajeev Kushwaha

📞 +91-7489502973 | ✉️ rajeev102003000@gmail.com | [LinkedIn](#) | [GitHub](#) | [PORTFOLIO](#) | 📍 Greater Noida, India

EDUCATION

G.L. Bajaj Institute of Technology & Management

B.Tech – Computer Science (AI & ML) | GPA: 7.48 | Media Head, Abhyudaya Club

Jan. 2023 – Dec. 2027

Greater Noida, India

Delhi Public School

Class XII – 81.2%

2020 – 2022

Delhi, India

Bhartiyam Vidhya Peeth

Class X – 97.3%

2019 – 2020

India

SKILLS

Languages : Python, C, C++, SQL

AI / ML : Machine Learning, Deep Learning, NLP, Generative AI, Agentic AI, AI Agents, Transformer & LLM Fine-Tuning (LoRA/QLoRA), LangChain, LangGraph, Pydantic, RAG Systems

Frameworks & Libraries : PyTorch, TensorFlow, Scikit-Learn, Pandas, FastAPI, Flask, Unsloth, BeautifulSoup

Developer Tools & Infra : Docker, GitHub, Git, MongoDB Atlas, Qdrant (Vector DB), Ollama, REST APIs

CS Fundamentals : DSA, DBMS, OOP, Operating Systems, Algorithms & System Design

AI-Assisted Dev : Claude, GitHub Copilot, ChatGPT for coding, debugging, automation, and rapid prototyping

PROJECTS

Advanced Agentic RAG Pipeline for AI Research Papers

[ragAPI](#) | [rag-advanced-research](#)

2025

- Built a production-style RAG system that parses academic PDFs via Nougat OCR / pymupdf4llm, preserving equations, tables, and section structure, and answers questions with cited sources
- Engineered hybrid dense-sparse retrieval combining BGE-M3/FastEmbed dense embeddings with BM25 sparse search, merged via Reciprocal Rank Fusion (RRF) over Qdrant's HNSW cosine index for stronger context quality
- Orchestrated a LangGraph ReAct agent loop with a retrieve → grade → generate pipeline and adaptive multi-LLM routing (Groq → Gemini → Ollama fallback based on available API keys)
- Served the pipeline via FastAPI with SSE token streaming and step-by-step pipeline tracking to a custom React chat interface; fully Dockerized and env-configurable with a toggleable local embedding model (BGE-small vs. MiniLM) to cut third-party embedding costs
- Stack**: LangGraph, Qdrant, FastEmbed (BGE-M3), Ollama, Groq, Gemini, Nougat OCR, FastAPI, React, Docker

AI Resume Checker & ATS Scanner | [Streamlit](#), [spaCy](#), [Sentence-Transformers](#), [Docker](#)

- Built an **ATS resume-scanner web app** that scores resume-to-job-description fit using a weighted blend of keyword matching (spaCy noun/proper-noun extraction) and **semantic similarity** via Sentence-Transformers (all-MiniLM-L6-v2).
- Designed a **composite scoring formula** (37% keyword match + 63% semantic similarity) that surfaces missing keywords and a 0–100% ATS match score; fully local, in-memory processing with no external API calls.
- Shipped with a **Streamlit UI** and **Docker** containerization for one-command deployment.
- Fine-tuned LLaMA (LoRA/QLoRA) via Unsloth for structured resume parsing; but it didn't performed expected

Text Summarizer | [Hugging Face Transformers](#), [FastAPI](#), [Docker](#) | [GitHub](#)

- Built an **end-to-end NLP text summarization pipeline** with a modular, configuration-driven architecture (YAML-based `config.yaml/params.yaml`) separating data ingestion, transformation, model training, and evaluation stages.

- Fine-tuned a **Hugging Face Transformer model** for abstractive summarization, evaluated with ROUGE/SacreBLEU metrics; achieved **88% accuracy**, outperforming traditional baselines like Naive Bayes.
- Exposed training and inference via **FastAPI** entry points; containerized with **Docker** for reproducible deployment.

Minesweeper RL Agent – Reinforcement Learning | *AMD Hackathon, IIT Delhi*

2024

- Trained a Minesweeper-playing agent using **GRPOTrainer** with custom reward shaping at the **AMD Hackathon, IIT Delhi**, running on a Jupyter-kernel virtual machine based on **Vertex AI Workbench**; evaluated exploration/exploitation trade-offs across training episodes.

CERTIFICATIONS & ACHIEVEMENTS

Generative AI and ChatGPT

Nov. 2025

GeeksforGeeks (Self-Paced)

Mastering Data Transformation Through NLP

Oct. 2025

GeeksforGeeks (Self-Paced)

Google Cloud Skill Badges

Silver League – 3272 pts

skills.google/public_profiles/rajeev-kushwaha

- * **7 badges earned:** Preparing for your Professional Data Engineer Journey, Build a Certification Study Guide: PMLE, MLOps with Agent Platform: Model Evaluation, Introduction to Large Language Models, MLOps with Agent Platform: Manage Features, MLOps for Generative AI, Introduction to Generative AI (*Oct. 2025 – Mar. 2026*).